

Kod_Ada_i_Grad

April 20, 2021

```
[1]: import pandas as pd
import numpy as np

from sklearn.datasets import load_boston
from sklearn.linear_model import LinearRegression, Lasso
from sklearn.metrics import mean_squared_error
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import FunctionTransformer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split

from matplotlib import pyplot as plt
import seaborn as sns
from scipy import stats
import matplotlib.image as mpimg

from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
```

```
[2]: df = pd.DataFrame(pd.read_json("https://api.apispreadsheets.com/api/dataset/
↳congressional-voting/",
                                orient = 'split'))

def encode(x):
    if x == "n":
        return -1
    if x == "?":
        return 0
    if x == "y":
        return 1
    if x == "republican":
        return -1
    if x == "democrat":
```

```

        return 1

df = df.drop_duplicates()
df = df.applymap( encode)

```

```

[3]: X = pd.DataFrame( df)
     y = pd.DataFrame( X["political_party"])
     X.drop( columns = ["political_party"], inplace=True )

```

```

[4]: X_train, X_test, y_train, y_test \
     = train_test_split(X, y, stratify = y, test_size=0.2, random_state=1)

X_train, X_val, y_train, y_val \
     = train_test_split(X_train, y_train, stratify = y_train, test_size=0.25,
     ↪random_state=1)

```

0.0.1 Baseline na podstawie physician_fee_freeze

```

[5]: def encode_base(x):
     if x== -1:
         return 1
     if x== 1:
         return -1
     if x == 0:
         return 0

```

```

[6]: print( 'Accuracy val: %.3f' % accuracy_score( y_val,
     ↪X_val["physician_fee_freeze"].map( encode_base )))
     print( 'Accuracy test: %.3f' % accuracy_score( y_test,
     ↪X_test["physician_fee_freeze"].map( encode_base )))

```

Accuracy val: 0.855
Accuracy test: 0.957

0.0.2 Adaboost

```

[7]: alf = AdaBoostClassifier(n_estimators=100, random_state=0, learning_rate=0.9) #
     ↪dobrze działające parametry
     alf.fit(X_train, y_train)

```

/home/kurowskik/anaconda3/lib/python3.8/site-packages/sklearn/utils/validation.py:73: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```

    return f(**kwargs)

```

```


[7]: AdaBoostClassifier(learning_rate=0.9, n_estimators=100, random_state=0)

```

```
[8]: y_val_hat = alf.predict(X_val)
      y_test_hat = alf.predict(X_test)
```

```
[9]: conf_matrix = confusion_matrix(y_true=y_val, y_pred=y_val_hat.round())
      fig, ax = plt.subplots(figsize=(5, 5))
      ax.matshow(conf_matrix, cmap=plt.cm.Oranges, alpha=0.3)
      for i in range(conf_matrix.shape[0]):
          for j in range(conf_matrix.shape[1]):
              ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center',
                      size='xx-large')

      plt.xlabel('Predictions', fontsize=18)
      plt.ylabel('Actuals', fontsize=18)
      plt.title('Confusion Matrix', fontsize=18)
      plt.show()
```



```
[10]: print('F1 Score: %.3f' % f1_score(y_test, y_test_hat))
       print('F1 Score: %.3f' % f1_score(y_val, y_val_hat))
```

```
print('Accuracy val: %.3f' % accuracy_score(y_val, y_val_hat))
print('Accuracy test: %.3f' % accuracy_score(y_test, y_test_hat))
```

F1 Score: 0.966

F1 Score: 0.978

Accuracy val: 0.971

Accuracy test: 0.957

0.0.3 GradBoost

```
[11]: from sklearn.ensemble import GradientBoostingClassifier
```

```
[12]: clf = GradientBoostingClassifier(n_estimators=90, learning_rate=0.8,
    max_depth=1, random_state=0).fit(X_train, y_train) # dobrze działające
    ↪parametry
    clf.score(X_test, y_test)
```


```
/home/kurowskik/anaconda3/lib/python3.8/site-
packages/sklearn/utils/validation.py:73: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return f(**kwargs)
```

```
[12]: 0.9565217391304348
```

```
[13]: y_val_hat = clf.predict(X_val)
    y_test_hat = clf.predict(X_test)
```

```
[14]: conf_matrix = confusion_matrix(y_true=y_val, y_pred=y_val_hat)
    fig, ax = plt.subplots(figsize=(5, 5))
    ax.matshow(conf_matrix, cmap=plt.cm.Oranges, alpha=0.3)
    for i in range(conf_matrix.shape[0]):
        for j in range(conf_matrix.shape[1]):
            ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center',
    ↪size='xx-large')

    plt.xlabel('Predictions', fontsize=18)
    plt.ylabel('Actuals', fontsize=18)
    plt.title('Confusion Matrix', fontsize=18)
    plt.show()
```



```
[15]: print('F1 Score val : %.3f' % f1_score(y_val, y_val_hat))
      print('F1 Score test: %.3f' % f1_score(y_test, y_test_hat))
      print('Accuracy val: %.3f' % accuracy_score(y_val, y_val_hat))
      print('Accuracy test: %.3f' % accuracy_score(y_test, y_test_hat))
```

```
F1 Score val : 0.978
F1 Score test: 0.966
Accuracy val: 0.971
Accuracy test: 0.957
```

0.0.4 HistGradBoost

```
[16]: from sklearn.experimental import enable_hist_gradient_boosting
      from sklearn.ensemble import HistGradientBoostingRegressor
```

```
[17]: est = HistGradientBoostingRegressor( l2_regularization= 2, learning_rate=0.7).
      ↪ fit(X_train, y_train)
      est.score(X_val, y_val)
```

/home/kurowskik/anaconda3/lib/python3.8/site-

```
packages/sklearn/utils/validation.py:73: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return f(**kwargs)
```


```
[17]: 0.5986765808634488
```

```
[18]: y_val_hat = est.predict(X_val).round()
      y_test_hat = est.predict(X_test).round()
```

```
[19]: conf_matrix = confusion_matrix(y_true=y_test, y_pred=np.int64(y_test_hat))

fig, ax = plt.subplots(figsize=(5, 5))
ax.matshow(conf_matrix, cmap=plt.cm.Oranges, alpha=0.3)
for i in range(conf_matrix.shape[0]):
    for j in range(conf_matrix.shape[1]):
        ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center',
        ↪size='xx-large')

plt.xlabel('Predictions', fontsize=18)
plt.ylabel('Actuals', fontsize=18)
plt.title('Confusion Matrix', fontsize=18)
plt.show()
```



```
[20]: print( accuracy_score(y_val, y_val_hat))
      print( accuracy_score(y_test, y_test_hat))
```

0.8695652173913043

0.9420289855072463

0.0.5 Test na danych tylko z dobrym rozdzielaniem

Sprawdziliśmy jak się zachowują wyniki, gdy wybieramy tylko te głosowania, które dobrze rozdzielają partie.

```
[21]: X2 = pd.DataFrame( X[ ["crime", "el_salvador_aid", "education_spending",
    ↪ "physician_fee_freeze", "adoption_of_the_budget_resolution"]])
```

```
[22]: X_train2, X_test2, y_train2, y_test2 \
      = train_test_split(X2, y, stratify = y, test_size=0.2, random_state=1)

X_train2, X_val2, y_train2, y_val2 \
      = train_test_split(X_train2, y_train2, stratify = y_train2, test_size=0.25,
    ↪ random_state=1)
```

```
[23]: alf2 = AdaBoostClassifier(n_estimators=100, random_state=0, learning_rate=0.9)
      alf2.fit(X_train2, y_train2)
```

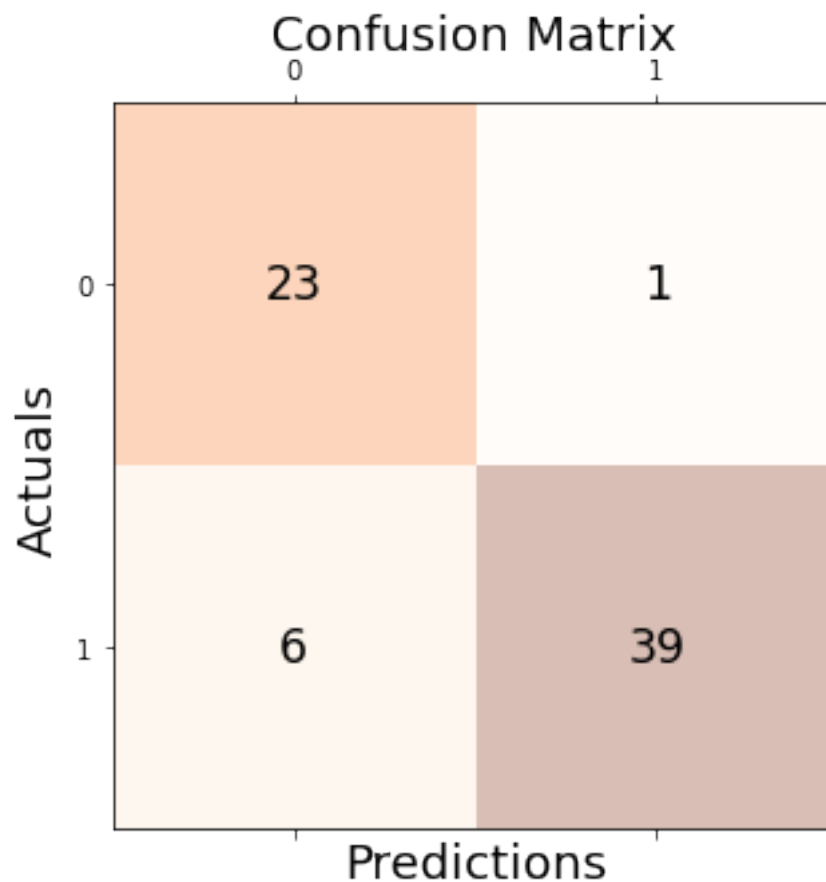
```
/home/kurowskik/anaconda3/lib/python3.8/site-
packages/sklearn/utils/validation.py:73: DataConversionWarning: A column-vector
y was passed when a 1d array was expected. Please change the shape of y to
(n_samples, ), for example using ravel().
    return f(**kwargs)
```

```
[23]: AdaBoostClassifier(learning_rate=0.9, n_estimators=100, random_state=0)
```

```
[24]: y_val_hat2 = alf2.predict(X_val2)
      y_test_hat2 = alf2.predict(X_test2)
```

```
[25]: conf_matrix = confusion_matrix(y_true=y_val2, y_pred=y_val_hat2.round())
      fig, ax = plt.subplots(figsize=(5, 5))
      ax.matshow(conf_matrix, cmap=plt.cm.Oranges, alpha=0.3)
      for i in range(conf_matrix.shape[0]):
          for j in range(conf_matrix.shape[1]):
              ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center',
                      size='xx-large')

      plt.xlabel('Predictions', fontsize=18)
      plt.ylabel('Actuals', fontsize=18)
      plt.title('Confusion Matrix', fontsize=18)
      plt.show()
```

```
[26]: print( accuracy_score(y_val2, y_val_hat2))
      print( accuracy_score(y_test2, y_test_hat2))
```

0.8985507246376812

0.9565217391304348

wyniki są nieco gorsze od Ada czy Grad, ale wciąż nieco lepsze od Baseline'u.

0.0.6 Usuwanie danych

Sprawdziliśmy, czy przy usunięciu słabo rozdziłających głosowań i 'outlierowych' kongresmenów uzyskamy lepsze rezultaty - okazało się, że nie.

```
[27]: X_count = np.apply_along_axis(sum, 1, abs(X))
      indexes = df[X_count < 11].index
      indexes = np.array(indexes)
      print(df.shape)
      X_dropped = df.drop(indexes, axis=0)

      print(X_dropped.shape)
```

(342, 17)

(333, 17)

```
[28]: col_ls = ['water_project_cost_sharing', 'immigration',  
             ↪ 'export_administration_act_south_africa', ]  
X_dropped = X_dropped.drop(col_ls, axis=1, errors='ignore')  
print(X_dropped.shape)
```

(333, 14)

```
[29]: y_dropped = X_dropped[['political_party']]  
X_dropped.drop(['political_party'], axis=1, errors='ignore', inplace=True)
```

```
[30]: X_dropped_train, X_dropped_test, y_dropped_train, y_dropped_test \  
      = train_test_split(X_dropped, y_dropped, stratify = y_dropped, test_size=0.  
      ↪ 2, random_state=1)  
  
X_dropped_train, X_dropped_val, y_dropped_train, y_dropped_val \  
      = train_test_split(X_dropped_train, y_dropped_train, stratify =  
      ↪ y_dropped_train, test_size=0.25, random_state=1)
```

```
[31]: alf_dropped = AdaBoostClassifier(n_estimators=100, random_state=0,  
      ↪ learning_rate=0.9)  
alf_dropped.fit(X_dropped_train, y_dropped_train)
```

```
/home/kurowskik/anaconda3/lib/python3.8/site-  
packages/sklearn/utils/validation.py:73: DataConversionWarning: A column-vector  
y was passed when a 1d array was expected. Please change the shape of y to  
(n_samples, ), for example using ravel().  
    return f(**kwargs)
```

```
[31]: AdaBoostClassifier(learning_rate=0.9, n_estimators=100, random_state=0)
```

```
[32]: y_dropped_val_hat = alf_dropped.predict(X_dropped_val)  
y_dropped_test_hat = alf_dropped.predict(X_dropped_test)
```

```
[33]: print( accuracy_score(y_dropped_val, y_dropped_val_hat))  
print( accuracy_score(y_dropped_test, y_dropped_test_hat))
```

0.9253731343283582

0.9253731343283582