

WUM_P2

April 20, 2021

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
```

```
[21]: df = pd.read_csv('congressional_voting_dataset.csv')
def encode(x):
    if x == "n":
        return -1
    if x == "?":
        return 0
    if x == "y":
        return 1
    if x == "republican":
        return -1
    if x == "democrat":
        return 1

df = df.drop_duplicates()
df = df.applymap(encode)

X = pd.DataFrame(df)
y = pd.DataFrame( X["political_party"])
X.drop( columns = ["political_party"], inplace=True )

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y,
↳test_size=0.2, random_state=1)

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, stratify =
↳y_train, test_size=0.25, random_state=1)
```

1 Modelowanie

Sprawdzam różne modele oraz tuninguje je różnymi metodami, aby uzyskać jak najwyższe accuracy. Jest to dobra miara do naszego zbioru ponieważ obie wartości targetu są dla nas tak samo istotne więc, ta miara sprowadza się do procentowej ilości poprawnych trafień.

```
[3]: from random import randint
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
from sklearn.metrics import confusion_matrix

from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import AdaBoostClassifier, RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.neighbors import KNeighborsClassifier
```

Model drzewa Wytrenowałem model pojedynczego drzewa tuningując go za pomocą GridSearch. Uzyskałem 91% w accuracy. Traktuję to jako baseline i teraz spróbuję znaleźć lepszy model.

```
[14]: tree = DecisionTreeClassifier()

param_grid = {"max_depth": range(3,16,3),
              "max_features": range(1, 40,5),
              "min_samples_leaf": range(1, 40,5),
              "criterion": ["gini", "entropy"]}

tree_Gcv = GridSearchCV(tree, param_grid, cv = 5)
tree_Gcv.fit(X_train, y_train)
print('Accuracy drzewa test: ' + str(round(tree_Gcv.score(X_test, y_test),3)))
print('Accuracy drzewa val: ' + str(round(tree_Gcv.score(X_val, y_val),3)))
```

Accuracy drzewa test: 0.942

Accuracy drzewa val: 0.87

```
[15]: plt.figure(figsize=(20,20))
plot_tree(tree_Gcv.best_estimator_, feature_names=X_train.columns.values,
          ↪fontsize=8, rounded=True, )
plt.show()
```



```

# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}

rf = RandomForestClassifier()
rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_iter = 100, cv = 3, verbose=2, random_state=42, n_jobs = -1)
rf_random.fit(X_train, y_train)
print('Accuracy Random Forest test: ' + str(round(rf_random.score(X_test, y_test), 3)))
print('Accuracy Random Forest val: ' + str(round(rf_random.score(X_val, y_val), 3)))

```

Fitting 3 folds for each of 100 candidates, totalling 300 fits
Accuracy Random Forest test: 0.971
Accuracy Random Forest val: 0.899

[25]: random_grid

```

[25]: {'n_estimators': [200, 233, 266, 300, 333, 366, 400, 433, 466, 500],
      'max_features': ['auto', 'sqrt'],
      'max_depth': [10, 11, 12, 13, 14, 16, None],
      'min_samples_split': [2, 5, 10],
      'min_samples_leaf': [1, 2, 4, 8, 12],
      'bootstrap': [True, False]}

```

1.1 Logistic Regression

```

[200]: lr = LogisticRegression()
grid = {'penalty' : ['l1', 'l2'], 'C' : np.logspace(-4, 4, 20)}
lr_Gcv = GridSearchCV(lr, param_grid=grid, cv=10)
lr_Gcv.fit(X_train, y_train)
print('Accuracy Logistic Regression test: ' + str(round(lr_Gcv.score(X_test, y_test), 3)))
print('Accuracy Logistic Regression val: ' + str(round(lr_Gcv.score(X_val, y_val), 3)))

```

Accuracy Logistic Regression test: 0.971
Accuracy Logistic Regression val: 0.913

1.2 Bayes

```
[201]: bayes = GaussianNB()
grid = {'var_smoothing': np.logspace(0,-9, num=100)}
bye_Gsv = GridSearchCV(bayes, param_grid=grid, cv = 10)
bye_Gsv.fit(X_train, y_train)
bye_Gsv.score(X_val, y_val)
print('Accuracy Bayes test: ' + str(round(bye_Gsv.score(X_test, y_test), 3)))
print('Accuracy Bayes val: ' + str(round(bye_Gsv.score(X_val, y_val), 3)))
```

Accuracy Bayes test: 0.957

Accuracy Bayes val: 0.884

1.3 KNN

```
[23]: knn = KNeighborsClassifier()
grid = {
    'n_neighbors': np.arange(1, 50),
    'weights': ['uniform', 'distance'],
    'metric': ['minkowski', 'manhattan']
}
knn_Gsv = GridSearchCV(knn, param_grid=grid, cv = 10)
knn_Gsv.fit(X_train, y_train)
print('Accuracy KNN test: ' + str(round(knn_Gsv.score(X_test, y_test), 3)))
print('Accuracy KNN val: ' + str(round(knn_Gsv.score(X_val, y_val), 3)))
```

Accuracy KNN test: 0.957

Accuracy KNN val: 0.884

1.4 AdaBoost

```
[203]: ada.get_params()
```

```
[203]: {'algorithm': 'SAMME.R',
      'base_estimator': None,
      'learning_rate': 1.0,
      'n_estimators': 50,
      'random_state': None}
```

```
[204]: ada = AdaBoostClassifier()
grid = {
    'n_estimators' : [50, 100, 200],
    'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.5, 0.7, 0.9]
}
ada_Gcv = GridSearchCV(ada, param_grid=grid, cv = 10)
ada_Gcv.fit(X_train, y_train)
print('Accuracy AdaBoost: ' + str(round(ada_Gcv.score(X_test, y_test), 3)))
```

Accuracy AdaBoost: 0.971

```
[208]: print('Accuracy AdaBoost: ' + str(round(ada_Gcv.score(X_val, y_val), 3)))
```

Accuracy AdaBoost: 0.884

```
[205]: ada2 = AdaBoostClassifier(n_estimators=100, random_state=0, learning_rate=0.9)
ada2.fit(X_train, y_train)
ada2.score(X_val, y_val)
```

```
[205]: 0.9710144927536232
```

```
[206]: df = pd.read_csv('congressional_voting_dataset.csv')
def encode(x):
    if x == "n":
        return -1
    if x == "?":
        return 0
    if x == "y":
        return 1
    if x == "republican":
        return -1
    if x == "democrat":
        return 1

df = df.drop_duplicates()
df = df.applymap(encode)

X = pd.DataFrame(df)
y = pd.DataFrame(X["political_party"])
X.drop( columns = ["political_party"], inplace=True )

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y,
↳test_size=0.2, random_state=1)

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, stratify =
↳y_train, test_size=0.25, random_state=1)
```

```
[207]: alf = AdaBoostClassifier(n_estimators=100, random_state=0, learning_rate=0.9)
alf.fit(X_train, y_train)
alf.score(X_val, y_val)
```

```
[207]: 0.9710144927536232
```

```
[214]: def c_matrix(model, X, y):
        y_val_hat = model.predict(X)
```

```

conf_matrix = confusion_matrix(y_true=y, y_pred=y_val_hat.round())
fig, ax = plt.subplots(figsize=(5, 5))
ax.matshow(conf_matrix, cmap=plt.cm.Oranges, alpha=0.3)
for i in range(conf_matrix.shape[0]):
    for j in range(conf_matrix.shape[1]):
        ax.text(x=j, y=i, s=conf_matrix[i, j], va='center', ha='center',
        ↪size='xx-large')

plt.xlabel('Predictions', fontsize=18)
plt.ylabel('Actuals', fontsize=18)
plt.title('Confusion Matrix', fontsize=18)
plt.show()

```

```
[215]: c_matrix(lr_Gcv, X_test, y_test)
```

